



EXPERIMENT - 5

Student Name: RAVI
Branch: CSE
Semester: 5th
Subject Name: ADBMS

UID: 23BCS10340
Section/Group: KRG 3-A
Date of Performance: 25/09/2025
Subject Code: 23CSP-333

1. Aim:

1. i. Create a large dataset:
 - a. Create a table names transaction_data (id , value) with 1 million records.
 - b. take id 1 and 2, and for each id, generate 1 million records in value column
 - c. Use Generate_series () and random() to populate the data.
- ii. Create a normal view and materialized view to for sales_summary, which includes total_quantity_sold, total_sales, and total_orders with aggregation.
- iii. Compare the performance and execution time of both.

2. The company **TechMart Solutions** stores all sales transactions in a central database

A new reporting team has been formed to analyze sales but **they should not have direct access to the base tables for security reasons.**

The database administrator has decided to:

- i. Create **restricted views** to display only summarized, non-sensitive data.
- ii. Assign access to these views to specific users using **DCL commands** (GRANT, REVOKE).

2. Objective:

- To create and populate a large dataset using SQL functions like generate_series() and random() for performance testing.
- To design normal views and materialized views for summarizing sales data, and to compare their performance in terms of execution time.
- To implement a secure reporting mechanism by restricting direct access to base tables and providing summarized, non-sensitive data through views.
- To apply Data Control Language (DCL) commands such as GRANT and REVOKE for assigning controlled access rights to specific users.
- To demonstrate how views and materialized views can improve query efficiency, security, and reporting in a sales transaction database.

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

3. Code:

1. Performance Benchmarking (Medium Level)

```
CREATE TABLE transaction_data (  
    id INT,  
    value INT  
);
```

```
INSERT INTO transaction_data (id, value)  
SELECT 1, (random() * 1000)  
FROM generate_series(1, 1000000);
```

```
INSERT INTO transaction_data (id, value)  
SELECT 2, (random() * 1000)  
FROM generate_series(1, 1000000);
```

```
SELECT COUNT(*) FROM transaction_data;
```

```
CREATE OR REPLACE VIEW sales_summary_view AS  
SELECT  
    id,  
    COUNT(*) AS total_orders,  
    SUM(value) AS total_sales,  
    AVG(value) AS avg_transaction  
FROM transaction_data  
GROUP BY id;
```

```
SELECT * FROM sales_summary_view;
```

```
EXPLAIN ANALYZE SELECT * FROM sales_summary_view;
```

```
CREATE MATERIALIZED VIEW sales_summary_mv AS  
SELECT  
    id,  
    COUNT(*) AS total_orders,  
    SUM(value) AS total_sales,  
    AVG(value) AS avg_transaction  
FROM transaction_data  
GROUP BY id;
```

```
SELECT * FROM sales_summary_mv;
```

```
EXPLAIN ANALYZE SELECT * FROM sales_summary_mv;
```

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

2. Securing Data Access with Views and Role-Based Permissions (Hard Level)

```
CREATE TABLE customer_master (  
    customer_id SERIAL PRIMARY KEY,  
    full_name VARCHAR(100)  
);
```

```
CREATE TABLE product_catalog (  
    product_id SERIAL PRIMARY KEY,  
    product_name VARCHAR(100),  
    unit_price NUMERIC(10,2)  
);
```

```
CREATE TABLE sales_orders (  
    order_id SERIAL PRIMARY KEY,  
    customer_id INT REFERENCES customer_master(customer_id),  
    product_id INT REFERENCES product_catalog(product_id),  
    order_date DATE,  
    quantity INT,  
    discount_percent NUMERIC(5,2)  
);
```

```
INSERT INTO customer_master (full_name) VALUES  
( 'Rvi Kumar'),  
( 'Tanya Verma'),  
( 'Alok Kumar'),  
( 'Neha Sharma');
```

```
INSERT INTO product_catalog (product_name, unit_price) VALUES  
( 'Laptop', 60000),  
( 'Keyboard', 1200),  
( 'Monitor', 15000),  
( 'Mouse', 800);
```

```
INSERT INTO sales_orders (customer_id, product_id, order_date, quantity,  
discount_percent) VALUES  
(1, 1, '2025-09-01', 1, 10),  
(2, 2, '2025-09-02', 2, 5),  
(3, 3, '2025-09-03', 1, 20),  
(4, 4, '2025-09-05', 3, 15);
```

DEPARTMENT OF

COMPUTER SCIENCE & ENGINEERING

CREATE OR REPLACE VIEW vw_order_summary AS

SELECT

O.order_id,

O.order_date,

P.product_name,

C.full_name,

(P.unit_price * O.quantity)

- ((P.unit_price * O.quantity) * O.discount_percent / 100) AS final_cost

FROM customer_master AS C

JOIN sales_orders AS O ON O.customer_id = C.customer_id

JOIN product_catalog AS P ON P.product_id = O.product_id;

CREATE ROLE rvi LOGIN PASSWORD '1234';

GRANT SELECT ON vw_order_summary TO rvi;

REVOKE SELECT ON vw_order_summary FROM rvi;

4. Output:

```
count
-----
2000000
(1 row)

CREATE VIEW

id | total_orders | total_sales | avg_transaction
---+-----+-----+-----
1 | 100000 | 499566284 | 499.566284000000000000000000000000
2 | 100000 | 499756817 | 499.756817000000000000000000000000
(2 rows)

QUERY PLAN
-----
Finalize GroupAggregate
  Group Key: transaction_data.id
  -> Sort
    Sort Key: transaction_data.id
    Sort Method: quicksort Memory: 25kB
    -> Gather
      Workers Planned: 2
      Workers Launched: 2
      -> Partial HashAggregate
        Group Key: transaction_data.id
        -> Parallel Seq Scan on transaction_data

Planning Time: 0.106 ms
Execution Time: 593.269 ms
(13 rows)

id | total_orders | total_sales | avg_transaction
---+-----+-----+-----
2 | 100000 | 499756817 | 499.756817000000000000000000000000
1 | 100000 | 499566284 | 499.566284000000000000000000000000
(2 rows)

QUERY PLAN
-----
Seq Scan on sales_summary_mv

Planning Time: 0.022 ms
Execution Time: 0.016 ms
(3 rows)

order_id | order_date | product_name | full_name | final_cost
---+-----+-----+-----+-----
1 | 2025-09-01 | Laptop | Rvi | 54000.00000000000000000000000000
2 | 2025-09-02 | Keyboard | Tanya Verma | 2280.0000000000000000000000000000
3 | 2025-09-03 | Monitor | Alok Kumar | 12000.0000000000000000000000000000
4 | 2025-09-05 | Mouse | Neha Sharma | 2040.0000000000000000000000000000
(4 rows)

ERROR: permission denied for view vw_order_summary
ERROR: permission denied for table customer_master
```



5. Learning Outcomes:

- Generate and manage large datasets using SQL functions like `generate_series()` and `random()`.
- Differentiate between normal views and materialized views with performance analysis.
- Secure data by allowing access only to summarized, non-sensitive information through views.
- Apply DCL commands (`GRANT`, `REVOKE`) to control user access and permissions.